


Summer School

Deep Learning

Prof. Dr. Klaus Dorer



Overview

- Neural Networks

- Introduction

- Model of a Neuron

- Perceptron

- Backpropagation Networks

- Convolutional Neural Networks

- Goals

- Know the elements of deep neural networks

- Have an estimation of their applicability

Why is deep learning so hot?

- More data
 - Youtube: 500 hours video per minute
 - Facebook: 300 million images per day
- More computing power
 - GPU acceleration
 - GPU/CPU Cluster
- ‚New‘ Deep Learning approaches
 - Deep Neural Networks
 - Convolutional Neural Networks
 - better activation functions, optimizer, initialization, ...
- Freely available frameworks
 - Keras, TensorFlow, Theano, DeepLearning4J, ...



Machine Learning before Deep Learning

Input






Feature Extraction
(Feature Engineering)

Klassifikation
Detektor


Result






SVM, (flat) neural network, ...




HMM, (flat) neural net, ...



Language recognition,



Clustering, HMM, ...



Machine translation, topic identification,


Machine Learning with Deep Learning

Input





Result




Deep
Network



„stapler“



„stapler“

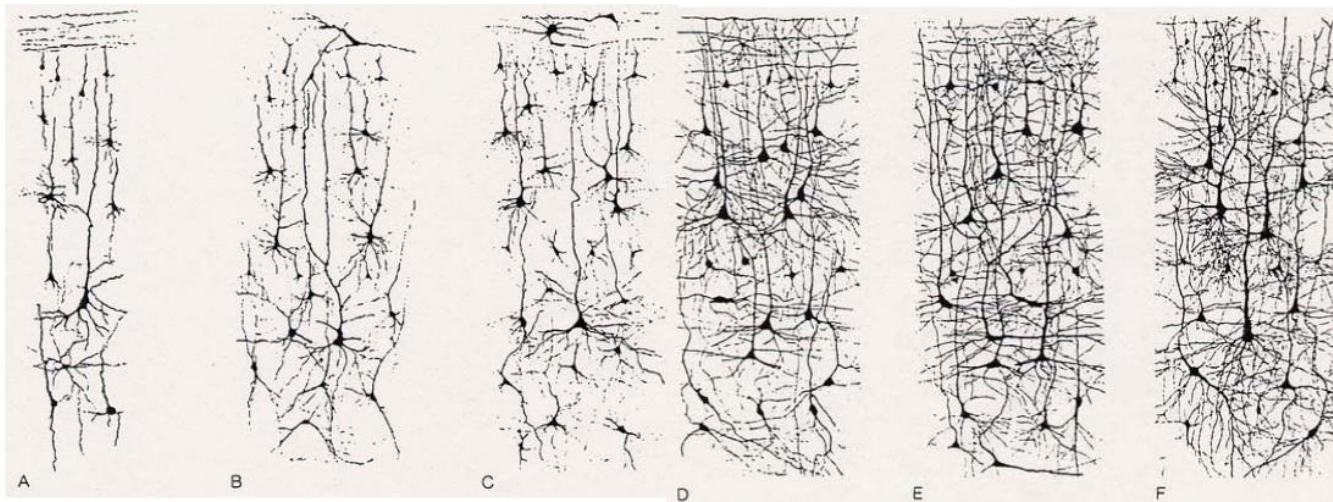


„Manual for a
stapler“

....

Human Brain


- Learning is performed by
 - Creating links between neurons
 - Reinforcing and weakening links between neurons



Die sechs Bilder vermitteln einen Eindruck von der Entwicklung des Gehirns von der Geburt bis zu einem Alter von zwei Jahren; zum Zeitpunkt der Geburt (A), nach einem Monat (B), nach drei (C), nach sechs (D), nach 15 (E) und nach 24 Monaten (F). Abgebildet ist ein Ausschnitt aus der Großhirnrinde in der Nähe des Broca Sprachareals.

Quelle: <http://nwg.glia.mdc-berlin.de/media/pdf/education/Legasthenie.pdf>


Human Brain



Deep Learning

Model of a Neuron


Prof. Dr. Klaus Dorer



Overview

- Neural Networks
 - Introduction
 - Model of a Neuron
 - Perceptron
 - Backpropagation Networks
 - Convolutional Neural Networks




Model of a Neuron



Model of a Neuron

Example: And Function

x ₁	x ₂	y ₁
0	0	0
0	1	0
1	0	0
1	1	1





$$o_j = \text{Step}_t\left(\sum_i w_{i,j} x_i\right) = \text{Step}_t(Wx)$$



Model of a Neuron

Or and Not Function

x_1	x_2	y_1
0	0	0
0	1	1
1	0	1
1	1	1







x_1	y_1
0	1
1	0







Model of a Neuron

Find suitable weights



Model of a Neuron



Find suitable weights



Model of a Neuron

Find suitable weights

■ Bias



$$o_j = \text{Step}_0\left(\sum_{i=1} w_{i,j} x_i\right)$$

Model of a Neuron

Find suitable weights

- Weights w are -1 and 1
- Loss function
 - Absolute error (L1 Norm)


$$l = \sum_i |y_i - o_i| = \|y - o\|_1$$

- Sum of squared error (L2 Norm)

$$l = \sum_i (y_i - o_i)^2 = \|y - o\|_2^2$$

- Gradient descent


x_1	x_2	y_1	o_1	l
0	0	0	0	
0	1	0	0	
1	0	0	1	
1	1	1	0	



Model of a Neuron

Summary

- Input function
- Bias
- Activation function




$$o_j = \text{Step}_0\left(\sum_{i=1} w_{i,j} x_i\right)$$

- Loss function

$$l = \|y - o\|_2^2$$


- Gradient descent



Deep Learning

Perceptron

Prof. Dr. Klaus Dorer




Overview

- Neural Networks
 - Introduction
 - Model of a Neuron
 - Perceptron
 - Backpropagation Networks
 - Convolutional Neural Networks

Perceptron

- Single layer feed forward network (1958 Frank Rosenblatt)



Perceptron

- Calculation of output

$$o_j = \text{Step}_0\left(\sum_{i=1} w_{i,j} x_i\right)$$

- Learning rule (delta rule)

$$w_{i,j} = w_{i,j} + \alpha \cdot x_i \cdot (y_j - o_j)$$

- Algorithm


Initialize weights randomly
do

 for each e in examples

 calculate output


 adjust weights

while (loss too high and other
 termination criteria not
 reached)




Perceptron

And Function



- Random start
 - $w_{x,0}$ [0.028; -0.258; -0.189]
 - 3 patterns wrong
- Train the four patterns
 - $w_{x,0}$ [-0.072; -0.158; -0.189]
 - 2 patterns wrong
- ...
 - 0 patterns wrong



Perceptron

Learning Rate and Gradient Descent

- Small learning rate



- Big learning rate




- Stochastic Gradient Descent
- Batch Learning

Perceptron

Real-valued Inputs and Outputs

$$o_j = \text{Step}_0\left(\sum_{i+1} w_{i,j} x_i\right)$$




- Winner takes all

- Softmax $\sigma(o_j) = \frac{e^{o_j}}{\sum_k e^{o_k}}$


Perceptron

Limitations


- Perceptrons can only represent linearly separable problems




AND



OR



XOR




Perceptron Summary

- Learning rule
 - Learning rate
 - Stochastic/Batch Gradient Descent
- Output function
 - Winner takes all
 - Softmax
- Problem linear separability

$$w_{i,j} = w_{i,j} + \alpha \cdot x_i \cdot (y_j - o_j)$$


$$\sigma(o_j) = \frac{e^{o_j}}{\sum_k e^{o_k}}$$



Deep Learning

Backpropagation Networks

Prof. Dr. Klaus Dorer



Overview

- Neural Networks
 - Introduction
 - Model of a Neuron
 - Perceptron
 - Backpropagation Networks
 - Convolutional Neural Networks

Backpropagation Networks


- To learn XOR oder similar problems we require
 - Multi-layer networks
 - A non-linear activation function
- Problem
 - How to adjust the weights of a hidden layer?
 - How can we propagate the error of the output back to previous layers?
- Solution
 - Backpropagation of Error
 - Bryson & Ho 1969, Rumelhart, Hinton & Williams 1986

Backpropagation Networks

Learning Rule

- Calculation of outputs

$$o_j = \sigma\left(\sum_{i+1} w_{i,j} x_i\right)$$



- Learning rule

- Weights to output

$$w_{j,k} = w_{j,k} + \alpha \cdot o_j \cdot \Delta_k \text{ mit } \Delta_k = \sigma'(in_k) \cdot (y_k - o_k)$$

- Weights to a hidden layer

$$w_{i,j} = w_{i,j} + \alpha \cdot o_i \cdot \Delta_j \text{ mit } \Delta_j = \sigma'(in_j) \cdot \sum_k w_{j,k} \Delta_k$$

Backpropagation Networks

Algorithm

```
initialize weights randomly
do
  for each e in examples
    calculate output (recall)
    calculate  $\Delta$  values for output units
    repeat for each layer(backward from the outputs)
      propagate  $\Delta$  values back to previous layer
      adjust weights
  while (loss too big and other stop criteria not reached)
```


Backpropagation Networks

Activation Function

- Sigmoid


- Function $\sigma(x) = \frac{1}{1 + e^{-x}}$

- Derivative $\sigma'(x) = \sigma(x)(1 - \sigma(x))$



- Tanh


$$\sigma(x) = \tanh(x)$$



- Relu

$$\sigma(x) = \max(0, x)$$



(Rectified linear unit)



Backpropagation Networks

Example: XOR Function


- A 3 – 3 – 1 network can learn XOR
 - Learning rate 0.9
- Learning rate 10



Backpropagation Networks

Example: classification


- What is a good size of the network?
 - To represent the training patterns?
 - Test of accuracy on training data
- Data
 - 864 patterns training
 - 864 patterns test
 - 6 attributes (4,4,4,3,3,3)
 - Class: yes/no
- Network
 - 21 input neurons
 - x hidden neurons
 - 2 output neurons



Backpropagation Networks

Generalization


- What is a good size of the network?
 - To properly predict unknown patterns?
 - Test of accuracy on test data



Backpropagation Networks

Learning Curve

- How many patterns are required for the network to generalize?
 - The more the better



Backpropagation Networks


Summary

- Learning rule


$$w_{i,j} = w_{i,j} + \alpha \cdot o_i \cdot \Delta_j \text{ mit } \Delta_j = \sigma'(in_j) \cdot \sum_k w_{j,k} \Delta_k$$

- Activation function


$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



- Network size




- Learning curve



Deep Learning

Convolutional Neural Networks

Prof. Dr. Klaus Dorer



Overview

- Neural Networks
 - Introduction
 - Model of a Neuron
 - Perceptron
 - Backpropagation Networks
 - Convolutional Neural Networks
 - Structure
 - Application
 - Deep Learning Frameworks

Convolutional Neural Networks

- Sweaty is a soccer robot
- It needs to see things on its camera images
- No matter where in the picture



- No matter how bright the light




- No matter what color, pattern




Convolutional Neural Networks

- Input is a 3D matrix
 - Here: 100x100 Pixel, 3 color chanel (Red, Green, Blue)




- Output is a feature-vector

- Ball
 - X-Line
 - Robot
 - L-Line
 - Obstacle



Convolutional Neural Networks


Example Architecture



Convolutional Neural Networks

Convolution Layer

- Filter (kernel) runs (convolves) over the input
- Calculates input for neuron in activation map
- Activation function here: ReLU
 - $Y = \text{ReLU}(Wx+b)$



$$\text{ReLU} \left(\begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 1 & 1 & 0 \\ 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 \\ 0 & 0 & 1 & 1 \end{bmatrix} + \begin{bmatrix} -1 \end{bmatrix} \right) = \begin{bmatrix} 0 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

- Repeated with n filters: n activation maps
- Padding: how do we deal with pixels at the border
- Strides: step size, > 1 means activation map shrinks

Convolutional Neural Networks

Pooling Layer

- Reduces the size of the activation map
- Pooling function
 - Usually Max-Pooling
 - Sometimes Avg-Pooling
- Example
 - 2x2 filter with 2,2 strides

2	3	5	7
6	4	3	2
1	2	3	2
0	1	1	0

Input



6	7
2	3

Output
Max- Pooling

Convolutional Neural Networks

Fully Connected Layer

- Already known: backpropagation layer
- Contains a majority of learnable weights
 - In our example
 - $2048 * 512 + 512 = 1.049.088$ weights
 - $512 * 8 + 8 = 4.104$ Gewichte
 - To compare with
 - $3*3*3*32 + 32 = 896$ weights in the first convolution layer
- On a pretrained network it may suffice to only learn these layers on new images



- Image tagging
 - A teacher has to assign the proper class to each image
 - In case of object localization also a bounding box is required
- Image learning
 - Present training image to network
 - Network calculates output and loss
 - Network performs gradient descent on loss function
- Accuracy is measured on test data
- Application
 - Stored network can be applied to live images
 - Fast recall

Convolutional Neural Networks


Overfitting

- Train many images
- What if we do not have too many?

- Augment existing images




- Create synthetic images

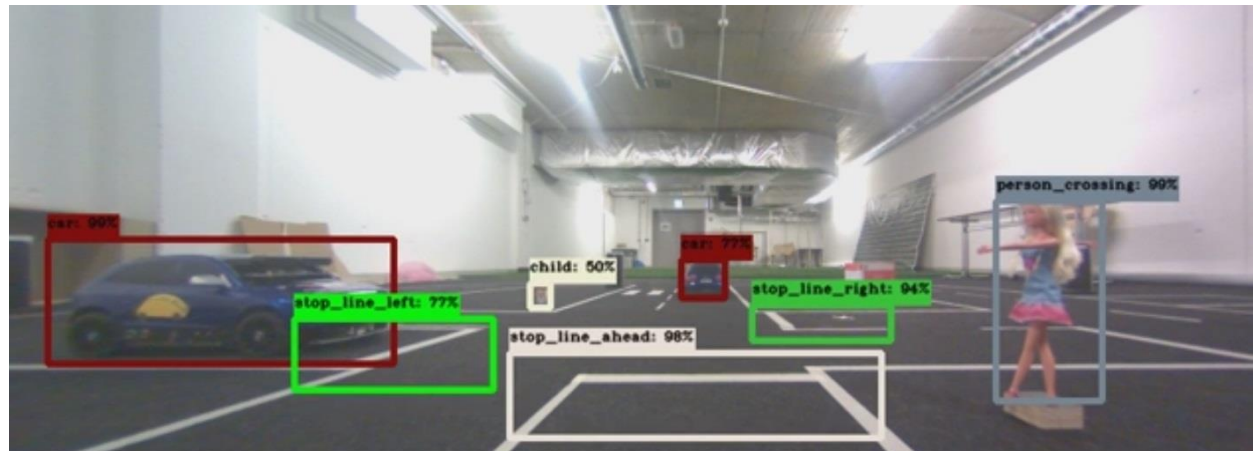


- Dropout
 - Do not use a fraction (e.g. 0.5) of random neurons during learning
 - Avoids single 'important' connections

- RoboCup




- AudiCup



Deep Learning

Applications: Google


- Input
 - 10 mio images (200 x 200 pixel)
- Learning
 - 1 Mio weights
 - 16.000 cores
 - 3 days



Deep Learning

Applications: NVIDIA

- Example NVIDIA
 - First layer detects lines and circles
 - Parts of a car
 - Cars
 - Car types
- Learning
 - Days
(on a GPU cluster)
- Recall
 - 2 Megapixel
 - 30 fps
 - 75 objects



Deep Learning

Applications: DeepMind

■ Deep Reinforcement Learning of computer games



■ Input



- 84*84*4 downsampled live video input

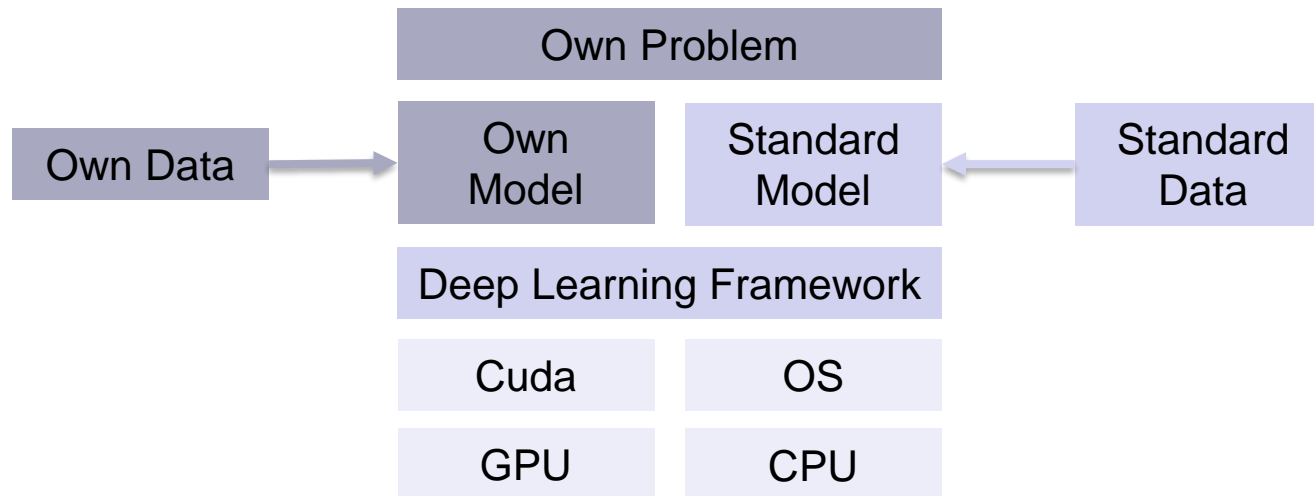
■ Network

- 4 layers, 2 convolutional (8x8, 4x4),
2 fully connected

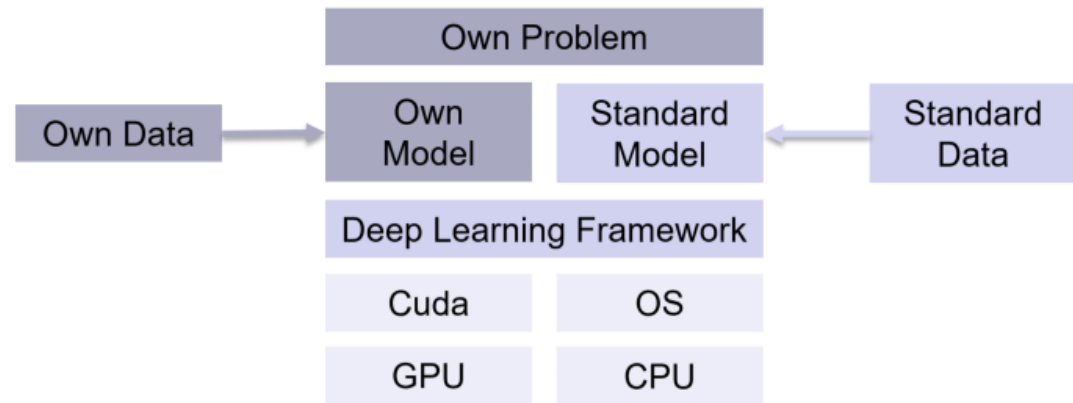
■ Result

- 4 of 7 games played better than human expert






- TensorFlow
 - Google Brain team
 - <https://www.tensorflow.org/>
- Torch, PyTorch
 - Communities
 - <http://pytorch.org/>
- Deeplearning4j
 - Skymind engineering team, Deeplearning4j community
 - <https://deeplearning4j.org/>
- Caffe
 - Berkeley Vision and Learning Center
 - <http://caffe.berkeleyvision.org/>
- Caffe2
 - Facebook
 - <https://research.fb.com/downloads/caffe2/>




- ImageNet
 - 14 Mio tagged images
 - 21.000 categories
 - <http://www.image-net.org>
- MNIST
 - 70.000 hand written digits
 - <http://yann.lecun.com/exdb/mnist/>
- COCO (common objects in context)
 - 200.000 tagged and segmented images
 - <http://cocodataset.org/#home>
- Music, faces, speech, texts, ...



Deep Learning

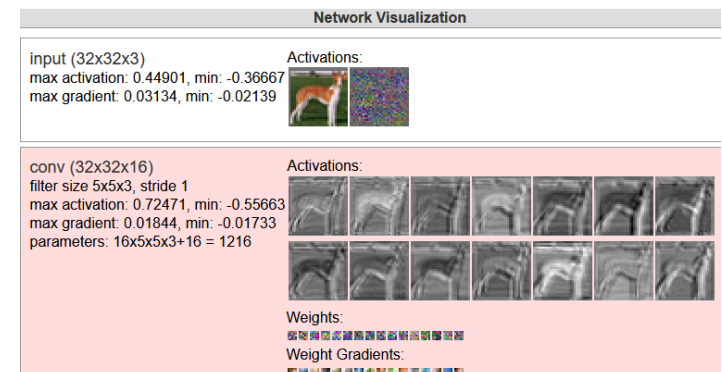

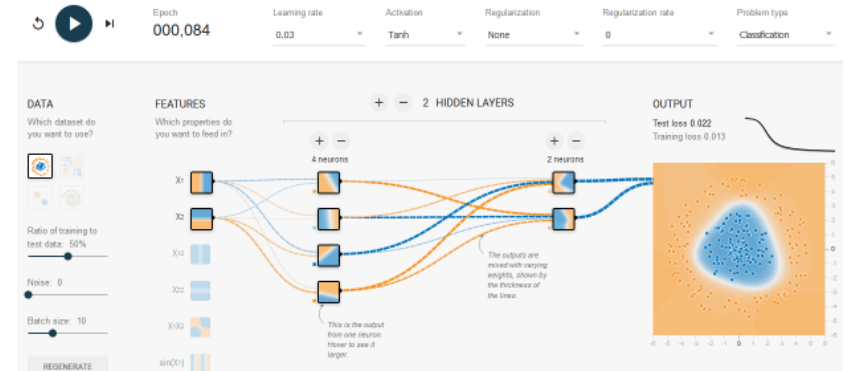
Standard Network Models

- LeNet-5 (1990)
 - 5 layers (4,1)
 - MNIST
- AlexNet (2012)
 - 8 layers (5,3)
 - ImageNet (16.4% Fehler)
- GoogLeNet (2014)
 - 22 layers
 - ImageNet (6.7%)
- ResNet-152 (2015)
 - 152 layers
 - ImageNet (3.6%), COCO



- VGGNet
- Mobilenet
- Inception
- ...

- Google Playground
 - <http://playground.tensorflow.org>
- 3D digit recognition (Adam Harley)
 - <http://scs.ryerson.ca/~aharley/vis/conv>
- ConvnetJS (Andrej Karpathy)
 - <https://cs.stanford.edu/people/karpathy/convnetjs>



- Deep Convolutional Neural Networks
 - Convolution layers
 - Max layers
 - Dense layers
- Applications
 - Image recognition, speech recognition, predictive maintenance, ...
- Many pretrained networks available for download



Europäischer Fonds für Regionale Entwicklung (EFRE)



Der Oberrhein wächst zusammen: mit jedem Projekt



UpperRhine 4.0